

CpE 313: Microprocessor Systems Design

Handout 03 Instruction Set Design

August 31, 2004

Shoukat Ali

shoukat@umr.edu



UNIVERSITY OF MISSOURI-ROLLA
The Name. The Degree. The Difference.

Credits

- for this and previous two lectures
 - Mazin Yousif, Intel
 - EECS Department at UC Berkeley

Issues in Chapter 2

- beginning chapter 2 today
- will learn
 - how to write a short program using stack, accumulator, and general-purpose-register machines
 - what a load-store architecture is
 - the advantages and disadvantages of these approaches
 - the instruction types, memory-addressing modes, control-flow instruction types, and operand sizes and types
 - what issues drive the choice of offset and immediate sizes in an instruction encoding
 - what an architecture should provide to help the compiler writer
 - what the different major memory segments are (stack, global-data, heap, and text)

3

Instruction Set Principles and Examples

- | | |
|-------|---|
| A n | Add the number in storage location n into the accumulator. |
| E n | If the number in the accumulator is greater than or equal to zero execute next the order which stands in storage location n ; otherwise proceed serially. |
| Z | Stop the machine and ring the warning bell. |

Wilkes and Renwick
*Selection from the List of 18 Machine
Instructions for the EDSAC (1949)*

4

Instruction Set Architecture: What Must be Specified?

- assume you just fetched a 16-bit long ALU instruction from memory and want to execute it
 - assume instruction = 1001 0011 0101 1110
- questions you should ask:
 - about instruction format or encoding
 - how does computer decode this binary sequence?
 - which bits specify the operation to be performed?
 - which bits specify the operands?
 - location of operands and result
 - how many operands must be *explicitly* specified?
 - zero in stack architecture
 - 1 in accumulator architecture
 - 2 or 3 in general purpose register architecture
 - are operands in memory or in processor storage?
 - how does the processor locate these operands?

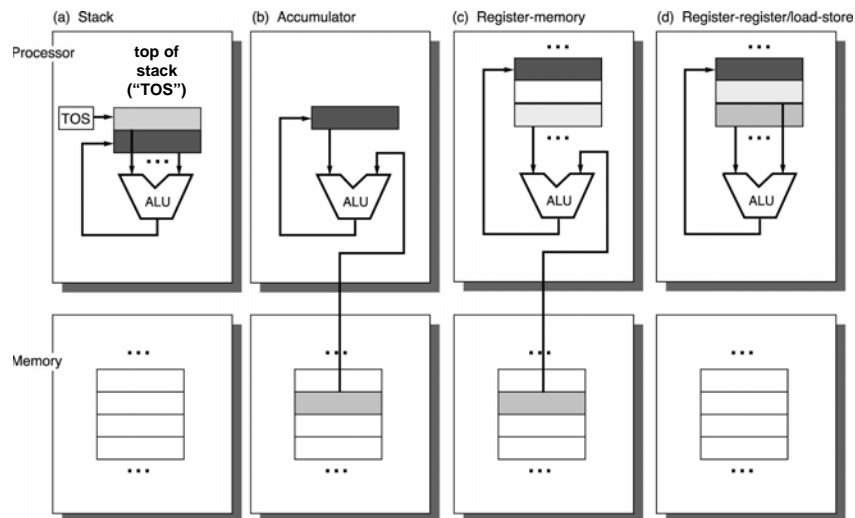
5

Classifying Instructions Sets – Location of Operands

- differentiating ISAs on the basis of location of **operands of an ALU operation** (add, sub, and, nand, etc)
- Stack Architectures
 - operands are at top of the stack
 - ALU operations do not have to specify operand locations
- Accumulator Architecture
 - one operand is stored in an accumulator
 - ALU ops do not have to specify this operand
 - other operand in memory
- General Purpose Register (GPR) Architecture
 - operands are in the general purpose register area or in memory
 - all operand locations have to be specified

6

Classifying Instructions Sets – Location of Operands



for memory-memory GPR arch, all operands are in memory

7

Doing $C = A + B$ in Three Basic ISA Classes

	stack	accumulator	register-register
$C = A + B$	$\text{TOS} \leftarrow \text{mem}(A)$ $\text{TOS} \leftarrow \text{mem}(B)$ add $\text{mem}(C) \leftarrow \text{TOS}$	$\text{ACC} \leftarrow \text{mem}(A)$ add $\text{mem}(B)$ $\text{mem}(C) \leftarrow \text{ACC}$	$\text{R1} \leftarrow \text{mem}(A)$ $\text{R2} \leftarrow \text{mem}(B)$ add $\text{R3}, \text{R1}, \text{R2}$ $\text{mem}(C) \leftarrow \text{R3}$

equivalent
representations

	stack	accumulator	register-register
$C = A + B$	push A push B add pop C	load A add B store C	load R1, A load R2, B add R3, R1, R2 store R3, C

register-register architecture also known as **load-store architecture**, implying that **only load and store instructions access memory**. Other instructions, e.g., ALU operations do NOT access memory.

8

A Bigger Example

	stack	accumulator	register-register
$A = B + C * D$ $E = B - C + D$	push B push C push D mult add pop A push D push C push B sub add pop E	load C mult D add B store A load B sub C add D store E	load , B load , C load , D mult R4,R2,R3 add R5, R4,R1 store R5, A sub R6, R1,R2 add R7, R6,R3 store R7, E

Question about load/store architecture: would it become harder to compile the C code if the processor had fewer registers? What if it had only 5 registers?

9

Register-Register Is Flexible!

stack	accumulator	register-register	register-register	register-register
push B push C push D mult add pop A push D push C push B sub add pop E	load C mult D add B store A load B sub C add D store E	load R1, B load R2, C load R3, D mult R4,R2,R3 add R5, R4,R1 store R5, A sub R6, R1,R2 add R7, R6,R3 store R7, E	load R3, D mult R4,R2,R3 add R5, R4,R1 store R5, A sub R6, R1,R2 add R7, R6,R3 store R7, E	load R3, D sub R6, R1,R2 add R7, R6,R3 store R7, E

for load/store arch, some instructions can be by compiler. This is important for caches and pipelining! We will see why later.

10

Code Size and CPU-to-Memory Traffic

- assume instruction operation codes are represented in 8 bits, memory addresses in 64 bit, register addresses in 6 bits, and data size is n bits (i.e., each access to memory moves n bits)
 - determine code size in bytes for each case
 - determine number of memory accesses in each case

stack			accumulator			register-register		
code	code size	data moved	code	code size	data moved	code	code size	data moved
push B	72	n	load C	72	n	load R1, B	78	n
push C	72	n	mult D	72	n	load R2, C	78	n
push D	72	n	add B	72	n	load R3, D	78	n
mult	8	0	store A	72	n	mult R4,R2,R3	26	0
add	8	0	load B	72	n	add R5, R4,R1	26	0
pop A	72	n	sub C	72	n	store R5, A	78	n
push D	72	n	add D	72	n	sub R6, R1,R2	26	0
push C	72	n	store E	72	n	add R7, R6,R3	26	0
push B	72	n				store R7, E	78	n
sub	8	0						
add	8	0						
pop E	72	n						
	608	8n		576	8n		494	5n

11

Pros And Cons

stack			accumulator			register-register		
code	code size	data moved	code	code size	data moved	code	code size	data moved
push B	72	n	load C	72	n	load R1, B	78	n
push C	72	n	mult D	72	n	load R2, C	78	n
push D	72	n	add B	72	n	load R3, D	78	n
mult	8	0	store A	72	n	mult R4,R2,R3	26	0
add	8	0	load B	72	n	add R5, R4,R1	26	0
pop A	72	n	sub C	72	n	store R5, A	78	n
push D	72	n	add D	72	n	sub R6, R1,R2	26	0
push C	72	n	store E	72	n	add R7, R6,R3	26	0
push B	72	n				store R7, E	78	n
sub	8	0						
add	8	0						
pop E	72	n						
	608	8n		576	8n		494	5n

← sum

- in load/store architecture, registers can be used to hold variables
 - this reduces traffic to memory chip
 - program speeds up because _____
 - code density improves because a register can be addressed with _____

12


How Many Registers Are Sufficient?

- how many registers are sufficient for a load/store architecture?

-

- program characteristics

```
load R1, B
load R2, C
load R3, D
mult R4, R2, R3
add R5, R4, R1
store R5, A
sub R6, R1, R2
add R7, R6, R3
store R7, E
```



```
load R1, B
load R2, C
load R3, D
mult R4, R2, R3
add R4, R4, R1
store R4, A
sub R1, R1, R2
add R1, R1, R3
store R1, E
```

reuse can hurt

```
load R2, A; load R3, B;
load R4, C; load R5, D
```

```
add R1, R2, R3
add R2, R4, R5 (reuse of R2)
```

versus

```
add R1, R2, R3
add R6, R4, R5 (no reuse)
```

without reuse "adds" are
"parallelizable" if there are two
adders (remember the ALU
enhancement from lecture 2)

- *Lots-of-registers* enable two important optimizations:
 - register allocation (more variables can be in registers)

-

13

Register-Register Types

- three operands
 - add R1, R2, R3
 - $R1 = R2 + R3$
- two operands
 - add R1, R2
 - $R1 = R1 + R2$

14

Five Minute Break

15

Memory Addressing

- issues
 - ■ how does the processor interpret a memory address?
 - how does an instruction specify a memory address?
 - “addressing modes”

16

Interpreting a Memory Address

- first issue: how does the processor interpret a memory address?
- ◦ usually a memory address is address of a byte
 - addresses 0, 1, 2, 3 point to successive bytes
 - addresses 0, 2, 4, 6 point to successive half-words
 - a half-word = 16 bits
 - addresses 0, 4, 8, 12 point to successive words
 - a word = 32 bits
- “processor wants to store a word B₀B₁B₂B₃ at address 4” means
processor will store bytes at addr 4, the next byte at
addr 5, ..., and the fourth byte at addr 7
 - which byte at address 4?
 - long political debate has raged over this issue; see “On Holy Wars and a Plea for Peace” at
<http://www.networksorcery.com/enp/ien/ien137.txt>

17

Big Endian or Little Endian

- big endian approach: store the most significant byte, B_0 , first
 - B_0 at addr 4, B_1 at addr 5, ..., B_3 at addr 7
 - IBM 360/370, Motorola 68k, Sparc, HP PA
- little endian approach: store the least significant byte, B_3 , first
 - B_3 at addr 4, B_2 at addr 5, ..., B_0 at addr 7
 - Intel 80x86, DEC VAX, DEC Alpha
- PowerPC and MIPS are bi-endian
 - can be programmed to any “end”
- - big endian: byte at address 4 is loaded as the most significant byte into the register
 - big endian: byte at address 4 is loaded as the least significant byte into the register

18


Object Alignment

- for objects that are larger than a byte in size, e.g., a word, double word, etc,
 - at which address do they start?
 - most computers require that objects start at an address which is a multiple of their size
 - e.g., word-size objects must start at addresses that are multiples of 4

Alignment: require that objects fall on address that is multiple of their size.

Aligned

Not Aligned





19

Why Enforce Object Alignment?

- DRAM/SRAM is so organized that a “memory access” reads/writes starting from a multiple of 4
- so asking for a word that begins at a multiple of 4 is okay

Asking for words beginning at 0 or 4 is OK. Asking for other words requires two reads (e.g., word starting at 2)



- so enforcing object alignment
 -
 -

20

Addressing Modes

- what we have done so far
 - discussed stack, accumulator, and GPR architectures
 - given an address for an object (word, double word etc.), we know which bytes to access in memory and how to determine which byte is MSB
- but how do we determine an address to start with?

21

How To Determine an Address: Addressing Modes

These are all the addressing modes used in recent computers

Addressing mode	Example	Meaning
Register	Add R4,R3	$R4 \leftarrow R4 + R3$
Immediate	Add R4,#3	$R4 \leftarrow R4 + 3$
Displacement	Add R4,100(R1)	$R4 \leftarrow R4 + \text{Mem}[100 + R1]$
Register indirect	Add R4,(R1)	$R4 \leftarrow R4 + \text{Mem}[R1]$
Indexed / Base	Add R3,(R1+R2)	$R3 \leftarrow R3 + \text{Mem}[R1 + R2]$
Direct or absolute	Add R1,(1001)	$R1 \leftarrow R1 + \text{Mem}[1001]$
Memory indirect	Add R1,@(R3)	$R1 \leftarrow R1 + \text{Mem}[\text{Mem}[R3]]$
Post-increment	Add R1,(R2)+	$R1 \leftarrow R1 + \text{Mem}[R2]; R2 \leftarrow R2 + d$
Pre-decrement	Add R1,-(R2)	$R2 \leftarrow R2 - d; R1 \leftarrow R1 + \text{Mem}[R2]$
Scaled	Add R1,100(R2)[R3]	$R1 \leftarrow R1 + \text{Mem}[100 + R2 + R3 * d]_{22}$

Addressing Mode Usage? (ignore register mode)

- 3 SPEC89 programs measured on a machine with all address modes (DEC VAX)
- results:

--- Displacement:	42% avg, 32% to 55%	↑75%
--- Immediate:	33% avg, 17% to 43%	↓85%
--- Register deferred (indirect):	13% avg, 3% to 24%	
--- Scaled:	7% avg, 0% to 16%	
--- Memory indirect:	3% avg, 1% to 6%	
--- Misc:	2% avg, 0% to 3%	

75% displacement & immediate

85% displacement, immediate & register indirect


23

Displacement Address Mode: Size?

Displacement

Add R4,100(R1)

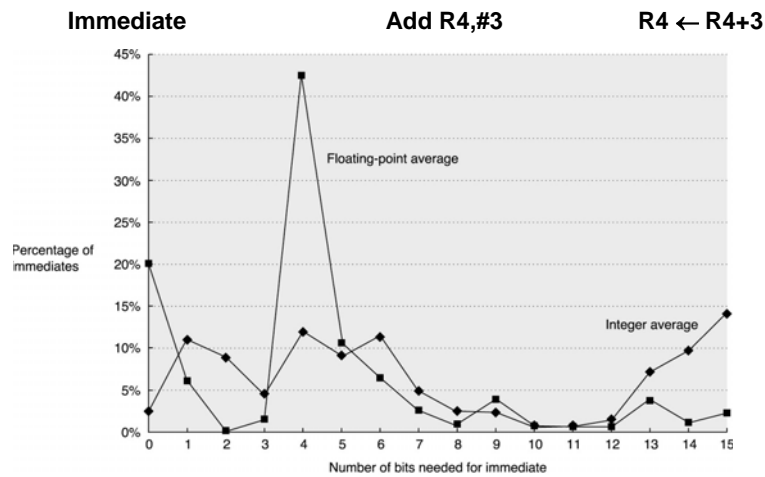
R4 ← R4+Mem[100+R1]



- ° Avg. of CINT2000 programs v. avg. of CFP2000 programs on Alpha
- ° only 1% of addresses needed more than 16-bits
- ° 12 - 16 bits of displacement needed

24

Immediate Address Mode: Size?



- 50% to 60% fit within 8 bits
- 75% to 80% fit within 16 bits

25

Addressing Mode Summary

- Data Addressing modes that are important:
 - Displacement, Immediate, Register Indirect
- Displacement size should be 12 to 16 bits
- Immediate size should be 8 to 16 bits

26

Type and Size of Operands

- what we have done so far
 - discussed stack, accumulator, and GPR architectures
 - given an address for an object (word, double word etc.), we know which bytes to access in memory and how to determine which byte is MSB
 - given an address field and an addressing mode, we know how to form the effective address
- now that we fetched the object from the memory, how do we know if we fetched an integer, a floating point number, a character, or some other data type?
 - btw, how many and what data types should a computer support?

27

Operand Types

Bit: 0, 1

Bit String: sequence of bits of a particular length

4 bits is a nibble
8 bits is a byte
16 bits is a half-word
32 bits is a word
64 bits is a double-word

Character:

ASCII 7 bit code
UNICODE 16 bit code

Decimal:

digits 0-9 encoded as 0000b thru 1001b
two decimal digits packed per 8 bit byte

Integers:

2's Complement

Floating Point:

Single Precision
Double Precision
Extended Precision


exponent
M x R
base
mantissa

IEEE standard 754 states how to specify a floating point number. See Appendix H for more details.

28

Operand Size Usage

Graph for sizes of objects referenced from memory
for SPEC CPU 2000 benchmarks



- a new 64-bit machine should support these data sizes and types:
 - 8-bit, 16-bit, 32-bit, 64-bit integers and
 - 32-bit and 64-bit IEEE 754 floating point numbers

29

What Kind of Operations Should a New Computer Support?

30

Typical Operations (little change since 1960)

Data Movement	Load (from memory) Store (to memory) memory-to-memory move register-to-register move input (from I/O device) output (to I/O device) push, pop (to/from stack)
Arithmetic	integer (binary + decimal) or FP Add, Subtract, Multiply, Divide
Shift	shift left/right, rotate left/right
Logical	not, and, or, set, clear
Control (Jump/Branch)	unconditional, conditional
Subroutine Linkage	call, return
Interrupt	trap, return
String	search, translate
Graphics (MMX)	parallel subword ops (4 16bit add)

31

Top 10 80x86 Instructions

rank	instruction	% of total executed
1	load	22%
2	conditional branch	20%
3	compare	16%
4	store	12%
5	add	8%
6	and	6%
7	sub	5%
8	move register-register	4%
9	call	1%
10	return	1%
	<hr/> Total	<hr/> 96%

◦ Simple instructions dominate instruction frequency

32


Operation Summary

Support these simple instructions, since they will dominate the number of instructions executed:

load,
store,
add,
subtract,
move register-register,
and,
~~shift,~~
compare equal, compare not equal,
branch,
jump,
call,
return;

33

Generic Examples of Instruction Format Widths



34

Variable Versus Fixed Length Instruction

- compile $E = A + B + C + D$
 - assume variables A,B,C,D have been loaded into registers 1 through 4
- | | |
|--|---|
| <ul style="list-style-type: none">■ variable length coding<ul style="list-style-type: none">■ can specify as many operands per instruction as one wants■ add R10,R1,R2,R3,R4■ assume opcode is 8 bits, reg code is 6 bits<ul style="list-style-type: none">■ code size = $8 + 30 = 38\text{bits}$ | <ul style="list-style-type: none">■ fixed length coding<ul style="list-style-type: none">■ can specify only two operands per instruction■ add R10,R1,R2,
add R10,R10,R3
add R10,R10,R4■ assume opcode is 8 bits, reg code is 6 bits<ul style="list-style-type: none">■ code size = $3(8 + 18) = 78\text{bits}$ |
|--|---|

35

Instruction Encoding Tradeoffs

- Variable width
 - Common instructions are short (1-2 bytes), less common or more complex instructions are long (>2 bytes)
 - (+) Very versatile,
 - (+) Small code size, memory is used efficiently
 - (-) Instruction words must be decoded before number of operands is known
- Fixed width
 - Typically 1 instruction per 32-bit word (DEC/Compaq Alpha is 2 instructions per word)
 - (+) Every instruction word is an instruction, Easier to fetch/decode
 - (-) Large code size, memory is used inefficiently

36

Instruction Formats

- If code size is more important, use variable length instructions
- If performance is more important, use fixed length instructions
- Recent embedded machines (ARM, MIPS) added optional mode to execute subset of 16-bit wide instructions (Thumb, MIPS16); per procedure decide performance or density

37

Instruction Mode Encoding

- Two possibilities for specifying addressing mode
 - Each operand has a “mode” field
 - Also called “address specifiers”
 - VAX, Motorola 68000
 - (+) Very versatile
 - (-) Encourages variable-width instructions (hard decode)
 - Opcode specifies addressing mode
 - Most RISCs (will see soon)
 - (+) Encourages fixed-width instructions (easy decode)
 - (+) “Natural” for a load/store ISA
 - (-) Limits what every instruction can do

38

Reduced Instruction Set Computing (RISC)

"The simplicity of all instructions allows a quick and simple evaluation of status to begin execution... Notice that this applies to all instructions. Adding complications to a special operation, therefore, degrades all the others."

-- J. E. Thornton, at the introduction of the CDC 6600, 1963

"I understand that in the laboratory developing this system there were only 34 people, including the janitor. .. Contrasting this modest effort with our own vast development activities, I fail to understand why we have lost our industry leadership by letting someone else offer the world's most powerful computer"

-- Thomas Watson, IBM president, in memo following 6600 announcement.

39

RISC Versus CISC

- | | |
|--|---|
| <ul style="list-style-type: none">■ small number of instruction sizes<ul style="list-style-type: none">■ preferably 1, e.g., MIPS■ 6600 had two■ only load/stores access memory<ul style="list-style-type: none">■ e.g. add instr cannot | <ul style="list-style-type: none">■ variable instruction size<ul style="list-style-type: none">■ 80x86 instruction could be anywhere from 1 to 17 bytes■ any instruction could access memory<ul style="list-style-type: none">■ add mem[A] mem[B] R2 |
|--|---|


40

Control Instructions

- instructions that change the flow of a program by explicitly changing the program counter
- types
 - conditional branches
 - if a certain condition is true, the program branches to an instruction called the “branch target” and proceeds from thereon
 - jump
 - program **unconditionally** jumps to an instruction called the “jump target” and proceeds from thereon
 - procedure call
 - program **unconditionally** jumps to another instruction, executes a set of instructions called a “procedure,” and then returns to the original program to resume the original flow
 - procedure return
 - the instruction that returns flow of control from a procedure to the original program

41

Conditional Branches Dominant



© 2003 Elsevier Science (USA). All rights reserved.

42

How to Specify Target Address

- displacement from current PC to target specified in the instruction
 - called “PC-relative” branches
 - requires fewer bits
 - experiments suggest 8 bits would work
 - gives the code “position independence”
- target specified in a register
 - example: jump (R1), i.e., jump to the address contained in R1
 - very useful if the target address is not known at compile-time as in a case statement

43

Conditional Branch: When to Compare

- possibility 1: combine compare calculation and branch instructions
 - $PC = PC + \text{displacement}$ if $R2 > R3$
- possibility 2: separate compare calculation and branch instructions
 - $R1 = R2 - R3$
 $PC = PC + \text{displacement}$ if $R1 > 0$
 - a variation on possibility 2: don't use up a general purpose register, instead provide and use a special register called CC (CC is like Program Status Word in 8051)
 - $R1 = R2 - R3$ (this op implicitly sets CC.NegFlag)
 $PC = PC + \text{displacement}$ if $CC.NegFlag == 0$

44

Procedure Calls and Returns

- we will postpone this until we have covered MIPS

45