

## Handout 19 TLBs and Virtual Caches

November 30, 2004

Shoukat Ali

shoukat@umr.edu

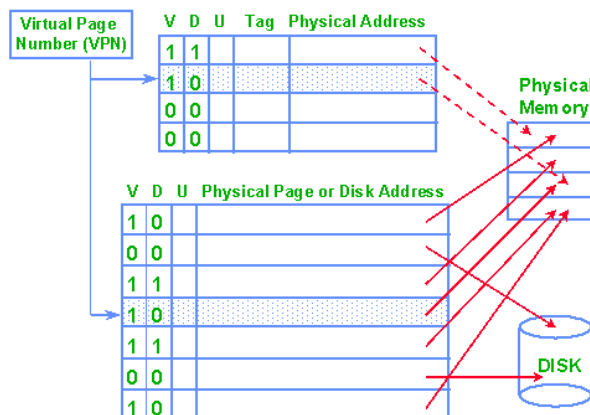


UNIVERSITY OF MISSOURI-ROLLA  
The Name. The Degree. The Difference.

1

### TLB Review

- each TLB entry holds a virtual address (called “tag”), the corresponding physical address, and some associated page protection and usage information




1. CPU VPN is compared against **ALL** stored VPNs (or tags)
2. if an entry matches, and valid bit is set, physical address is taken out of TLB

2

- a CPU routinely switches between programs (processes)
- process has its            page table
- when a new process is “switched in,” the new process’s page table must be used
- there is            TLB for all processes
- is it true that VPN→PPN mappings stored in TLB become meaningless at every context switch
  - yes
  -

33

- in some TLBs, each entry is augmented with a ASN field
  - ASN = address space identifier (aka PID = process ID)




1. *ASN*: VPN from CPU is compared against *ALL* stored *ASN*: VPN values
2. if an entry matches, valid bit is set, *and no protection bits are violated*, then physical address is taken out of TLB

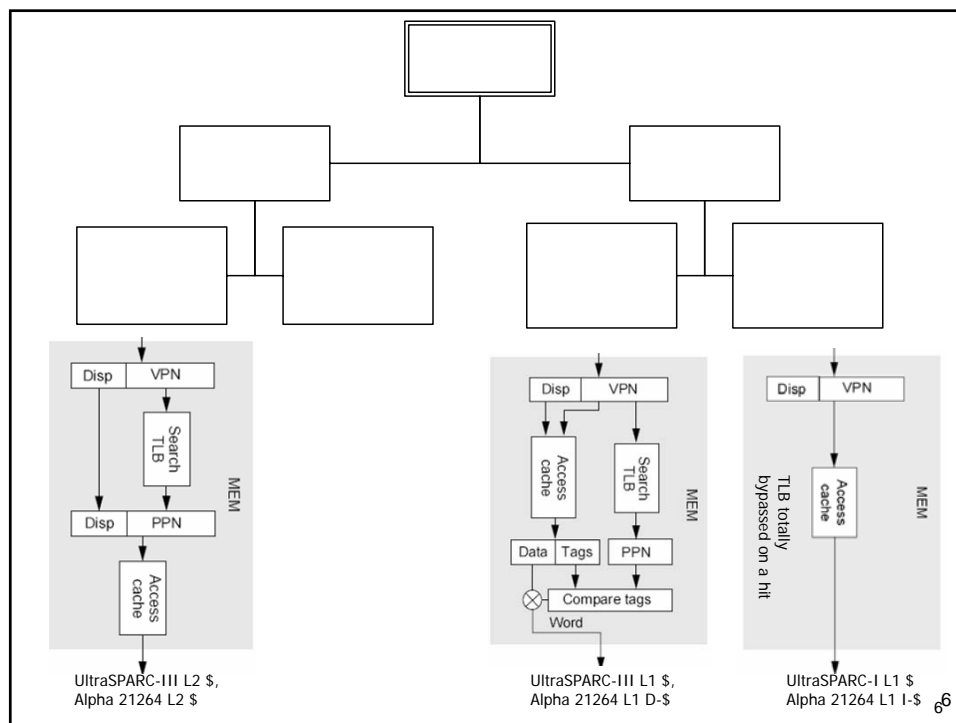
- such TLBs do not need to be flushed at every context switch
  - why?
- at what type of context switch, should this TLB be flushed?

44

## Which Cache Organization Can Bypass TLB?



5




## Is Bypassing TLB Fine?

- page table and TLB (which is just cached page table) are like
  - contains protection information that can be checked against
  - bypassing TLB means a memory request can bypass the police checkpoint; NOT GOOD
- *for a virtually indexed, virtually tagged cache, stored in cache are virtual and are compared with the virtual tag*
  - → TLB is totally bypassed on
  - page table is accessed
- so how can one re-design V/V caches so that protection information is checked for cache hits?

77

## Enforcing Protection Checks in V/V Caches

- how can one re-design V/V caches so that protection information is checked for cache hits?
- solution:
  - store protection info in both TLB and the cache
  - check the protection information at each access to V/V cache



88

## Should a V/V \$ be Flushed At Process Switch?

- a CPU routinely switches between programs (processes)
- process has its page table
- when a new process is “switched in,” the new process’s virtual addresses
- there is cache for all processes
- is it true that (virtual tag→block) mappings stored in a V/V cache become meaningless at every context switch?
  - a V/V cache must be incoming process fills up the cache as it proceeds

} just like in a TLB

99

## A Better Solution

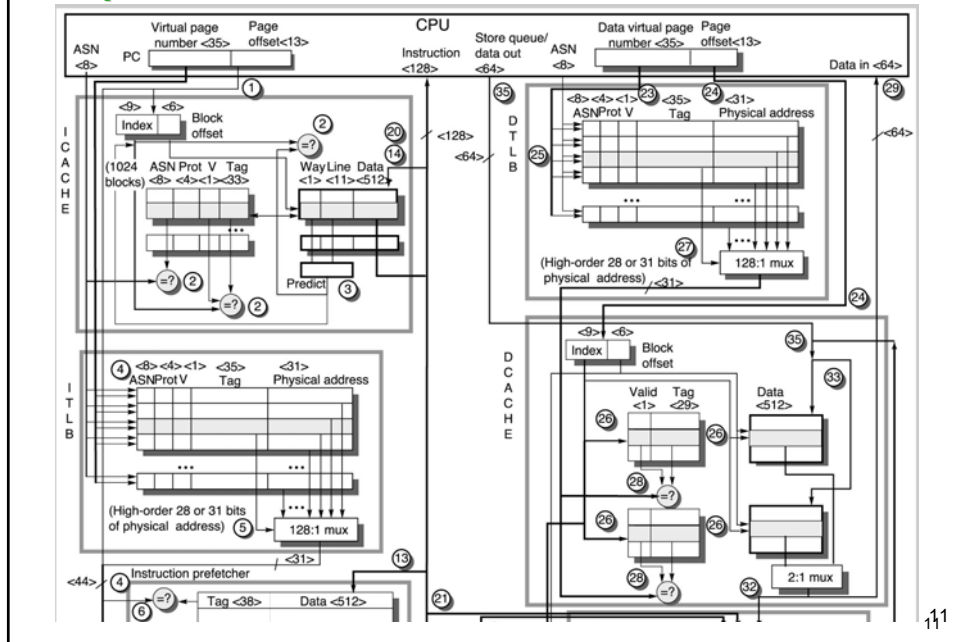
- augment V/V cache line with a PID (process ID) field
- the CPU always sends PID and VA to the cache
  - in effect, PID:VA make the true virtual address
- the search process changes in the following manner:
  - PID:VA from the CPU is compared against the stored PID:Tag values

| PID | Prot | V   | Tag | Data |
|-----|------|-----|-----|------|
|     |      |     |     |      |
|     |      |     |     |      |
|     |      |     |     |      |
| ... | ...  | ... | ... | ...  |
|     |      |     |     |      |

- what if a process dies, and its PID is re-cycled (i.e., assigned to a new born process)?
  - now the cache must be flushed

100

## Alpha 21264 Test: Which Cache is a V/V Cache?



## For Reference

Table 2. Cache organizations of some microprocessors.

| Microprocessor      | First-level instruction cache      | First-level data cache             | Second-level cache                         |
|---------------------|------------------------------------|------------------------------------|--|
| Intel Pentium       | 8 Kbytes; 2-way SA*; P/P           | 8 Kbytes; 2-way SA; P/P            | No support                                 |
| Intel PentiumPro    | 8 Kbytes; 2-way SA; P/P            | 8 Kbytes; 2-way SA; P/P            | 256 Kbytes; 4-way SA; P/P                  |
| DEC Alpha 21164     | 8 Kbytes; DM**; V/V                | 8 Kbytes; DM; P/P                  | 96 Kbytes; 3-way SA; P/P                   |
| PowerPC 604         | 16 Kbytes; 4-way SA; P/P           | 16 Kbytes; 4-way SA; P/P           | No support                                 |
| PowerPC 620         | 32 Kbytes; 8-way SA; P/P           | 32 Kbytes; 8-way SA; P/P           | 1 to 128 Mbytes; DM; P/P (external)        |
| H-P PA7100          | Up to 1 Mbyte; DM; V/P (external)  | Up to 2 Mbytes; DM; V/P (external) | No support                                 |
| H-P PA 8000         | Up to 4 Mbytes; DM; V/P (external) | Up to 4 Mbytes; DM; V/P (external) | No support                                 |
| Mips R4400          | 64 Kbytes; DM; V/P                 | 64 Kbytes; DM; V/P                 | 128 Kbytes to 4 Mbytes; P/P (external)     |
| Mips R8000          | 16 Kbytes; DM; V/V                 | 16 Kbytes; DM; V/P                 | 1 to 16 Mbytes; 4-way SA; P/P (external)   |
| Mips R10000         | 32 Kbytes; 2-way SA; V/P           | 32 Kbytes; 2-way SA; V/P           | 1/2 to 16 Mbytes; 2-way SA; P/P (external) |
| Sun Ultraspac I, II | 16 Kbytes; 2-way SA; P/P           | 16 Kbytes; DM; V/P                 | 512 Kbytes to 4 Mbytes; P/P (external)     |


\*Set associative

\*\*Direct mapped

Table taken from "VIRTUAL-ADDRESS CACHES, Part 2: Multiprocessor Issues," IEEE Micro, November/December 1997, pages 69—74.

## The Issue of Synonyms in VI Caches

- issue: programs sometimes use two or more different virtual addresses for the same physical address
  - these duplicate VA are called synonyms or aliases
  - a VI cache may keep two or more copies of one physical location
  - small problem: *cache pollution*



- big problem: what if two synonyms are in the cache, and one is modified by a program?

13

## Solution to Synonym Problem

- *synonym problem: the problem of ensuring that a block must never co-exist in the cache under different addresses*
- solution: before bringing in a new block in cache, search the VI cache for any synonyms if any are found
- how do we search a VI cache for synonyms of v1
  - V/P cache: for each cache block being examined, retrieve the stored tag, and see if it matches tag obtained by TLB translation of v1
  - V/V cache: for each cache block being examined, retrieve the stored tag, and see if it matches tag obtained by TLB translation of v1

14

## Two Implementations of the Solution

- how do we implement the above solution?
  - UltraSPARC – uses a DMC along with OS-implemented page coloring
  - Alpha 21264 – implemented by HW using antialiasing

15

## Key Observation About Synonyms - 1

- by definition, synonyms v1 and v2 are virtual addresses that translate to one physical address p
    1. above implies they have exact same value
    2. can only differ
- |     |             |    |
|-----|-------------|----|
| VPN | Page Offset |    |
| Tag | Index       | BO |
3. item 2 implies that indices for synonyms may differ only in superset bits (if there are any superset bits)
- grand conclusion: all synonyms of v1 have to be formed by v1's index

16



## Locating Synonyms in a Set Associative Cache

- assume memory system wants to locate any synonyms of v1
  - i.e., in which cache lines, synonyms of v1, if any, are residing?
- based on grand conclusion (slide 16), determine if the memory system should search the
  - entire cache
  - entire set indexed by v1
  - entire superset formed by v1's index
- to locate v1's synonyms,

177

### Example: Locating Synonyms in a Set Associative Cache

- assume a 4-way set associative cache, an index of 6 bits, and a page offset of 4 bits equal to `iiii`
- to locate `v1`'s synonyms, system needs to search the superset formed by `v1`'s index

|    | ASN  | Prot | V | Tag | Data |
|----|------|------|---|-----|------|
| 00 | iiii |      |   |     |      |
| 01 | iiii |      |   |     |      |
| 10 | iiii |      |   |     |      |
| 11 | iiii |      |   |     |      |
|    | ...  |      |   |     |      |

 ${}_{18}^{18}\text{O}$

## Alpha 21264's Solution to Synonym Problem

- *synonym problem: the problem of ensuring that multiple copies of a block must never co-exist in the cache under different addresses*
- assume
  - v1 and v2 are two synonyms
  - CPU sends a request for v2
- hardware searches the entire superset to see if tag for any block matches physical address of v2
  - if yes, (which will happen if v1 is in cache)
    - bring the block from main memory and store it in cache in v2's set
  - if no, (which will happen if v1 is not in cache)
    - bring the block from main memory and store it in cache in v2's set

19

## Synonym Alignment

- accessing entire superset in Alpha's case is easy because there are only eight blocks in the superset
- superset access gets ve the number of superset bits increases
- is there any way that HW does not have to search the entire superset?
- yes:
  - OS places the restriction that all synonyms must have the same superset bits
  - OS places the restriction that superset bits of the physical page are same as that of virtual page (aka: both pages have the same "color")
  - called "synonym alignment"
  - done in SunOS as part of a scheme called "page coloring"

20

### UltraSPARC-III's Solution to Synonym Problem

- with synonym alignment: *"HW searches only one set to see if tag for any block matches physical address of v2"*
- with DMC: a cache set contains only one block
- with synonym alignment + DMC,
  - CPU sends out a request for v2
  - v2 is brought in
  - if v1 is in cache, it is in the same block frame where v2 will be housed
    - v1 is kicked out to make space for v2
    - no need to search
- UltraSPARC-III uses a DMC along with synonym alignment to solve the synonym problem

21

### Cost of Aligning Synonyms

- page coloring imposes constraints on memory allocation
- when a new physical page is allocated on a page fault, the memory management algorithm must pick a page with the same color as the virtual color from the free list
- memory is not fully associative any more
  - page fault rate will likely increase

22