

CpE 313: Microprocessor Systems Design Fall 2004

Cache Coherence Handout 20

December 2, 2004

Shoukat Ali

shoukat@umr.edu



UNIVERSITY OF MISSOURI-ROLLA
The Name. The Degree. The Difference.

1

Dell PowerEdge 6600 Server – A Shared Memory System



list price for a 4-node
system = \$4,199

2

IBM eServer xSeries 440 – Another Shared Memory System



3

Sun Fire 12K Server – Yet Another Shared Memory System



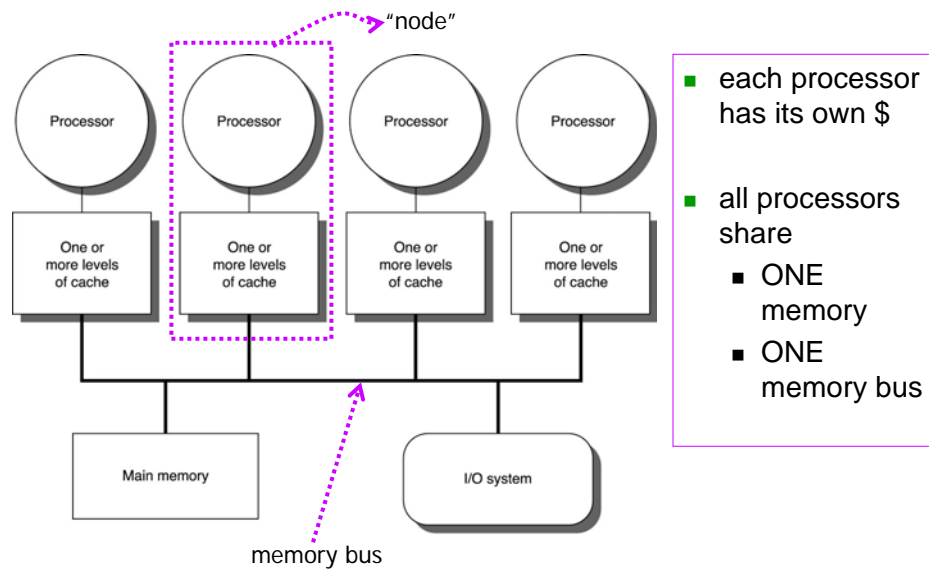
list price for:

a 4-node system
= \$354,330

a 36-node system
= \$1,382,530

4

What is a Shared Memory System?



5

Cache Coherence Problem in SM Multiprocessors

- only issue that we can investigate in this course
- will set one problem (up to 5%) on it in the final exam

6

Cache Coherence Problem in SM Multiprocessors

- example of the problem for a 2-node SM multiprocessor
- processor 1 reads location A
- processor 2 reads location A
- processor 1 writes location A
- now, processor 2 has stale data in its cache
- solution: if proc 1 wants to write location A, it must
 - either update copy of A in the \$ of proc 2
 - or invalidate copy of A in the \$ of proc 2
- “or” part of the solution is very similar to the way Alpha 21264 handles the synonym problem

7

Cache Coherence Protocols

- informal definition of cache coherence
 - a method of making memory accesses appear to be consistent despite the presence of caches
- cache coherence (CC) protocol: a protocol for maintaining cache coherence
- two types of CC protocols depending on how writes are handled
 - invalidate copies in other caches
 - invalidation protocols
 - update copies in other caches
 - update protocols

8

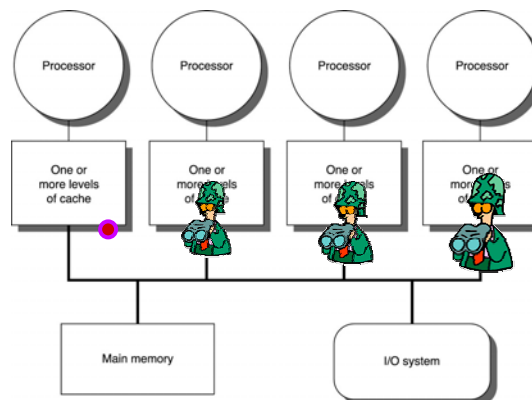
Invalidation-Based Cache Coherence

- basic idea
 - a block of data can be shared among multiple processors as long as no processor wants to write to it
 - if a processor wants to write a cache block, it must ensure that no other processor has a copy of this block
 - if yes, that copy must first be invalidated
- how does a processor know a data item is being shared?
 - have its cache snoop on other processors to see what blocks they are requesting from the memory

9

Snooping Caches

- a snoopy cache is like a snoopy next door neighbor who is always watching to see what you are doing
- a snoopy cache is always monitoring the shared memory bus to see what other processors are requesting from the memory



10

Illinois (or MESI) Protocol (Pentium, PowerPC, MIPS)

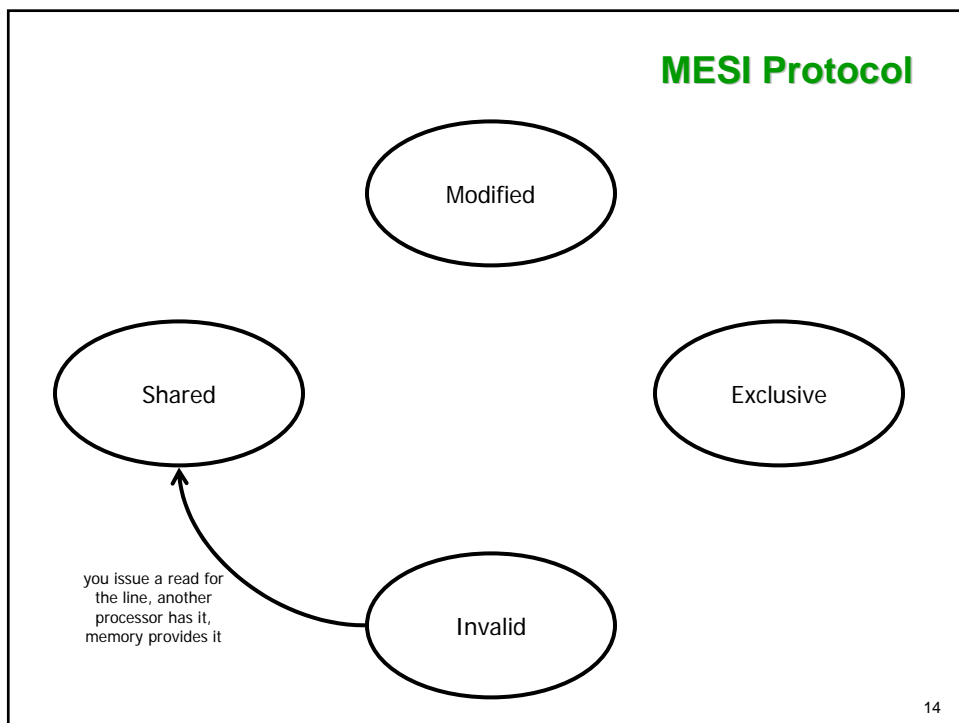
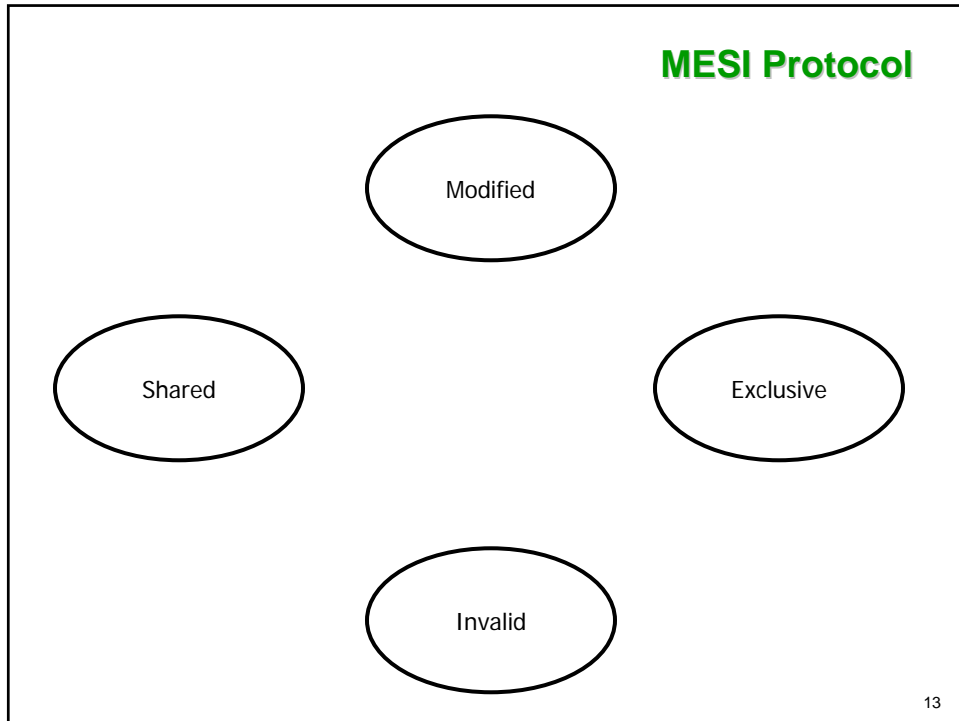
- in a processor's cache, each block can be in one of four states
 - modified (this processor has the only copy, the copy is “dirty,” i.e., it has been modified since fetched, memory is “stale”)
 - exclusive (this processor has the only copy, the copy is “clean,” i.e., it has not been written since fetched, memory is up-to-date)
 - shared (this processor and some other processors have identical copies of this block, memory is up-to-date)
 - invalid (this processor has no copy of the block)
- a cache records the states of its blocks
 - similar to how uniprocessor caches track valid/invalid

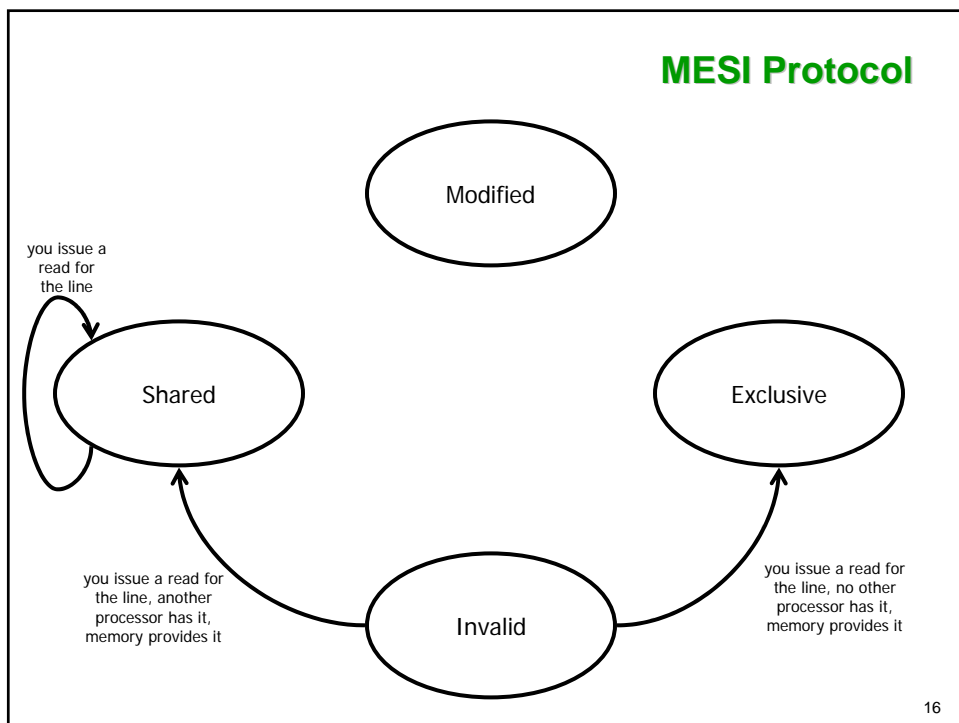
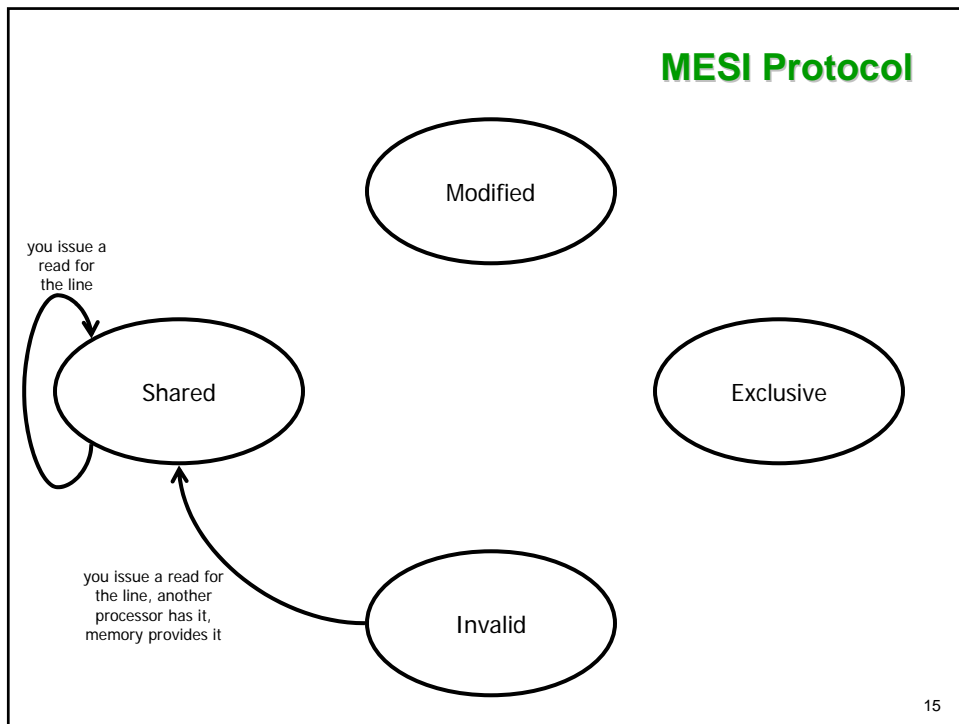
11

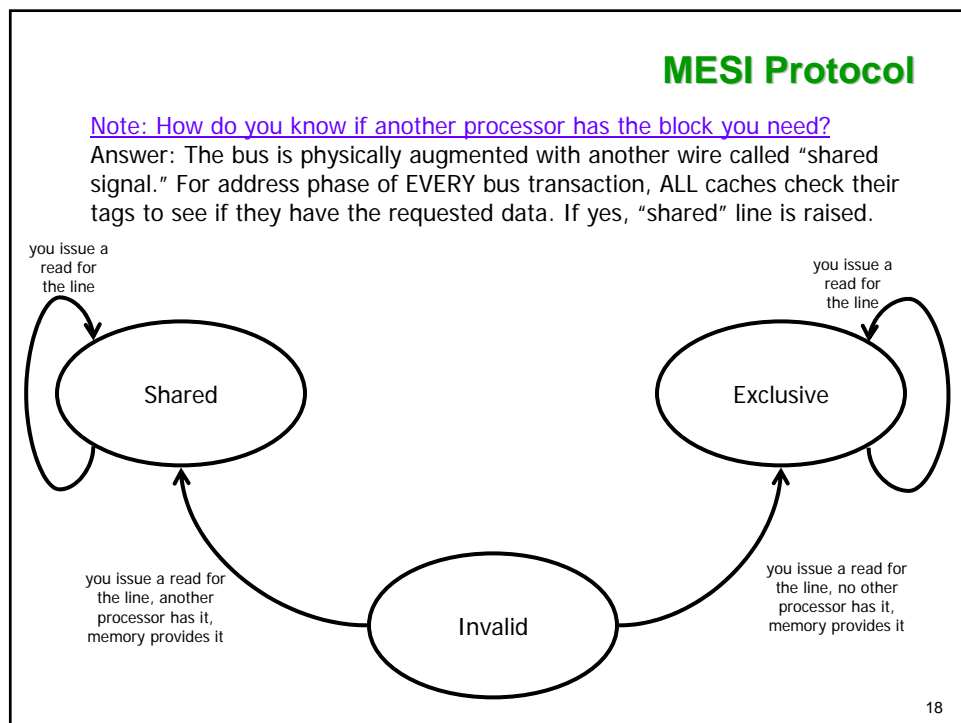
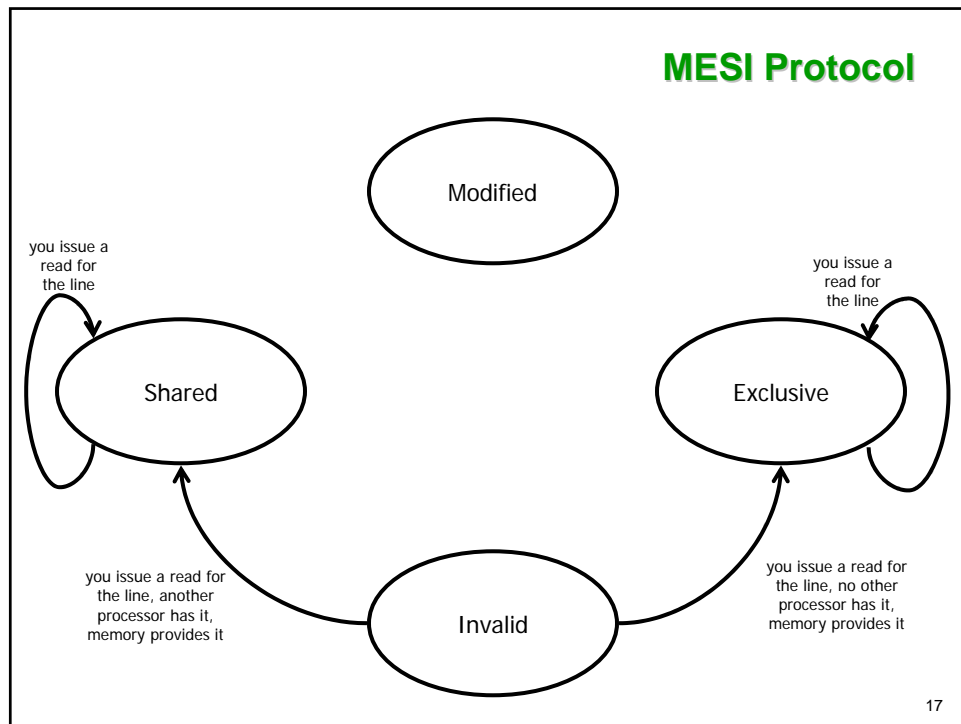
Interpreting the Illinois FSM

- FSM is drawn from the perspective of a given processor called “you”
- there are other processors as well
- “your” cache (as well as other caches) are monitoring the SM bus
 - your cache will either see a memory request from you or from another processor
 - your cache will compare the requested address with its tags
 - if address is found, it will take an action

12







Some Notation

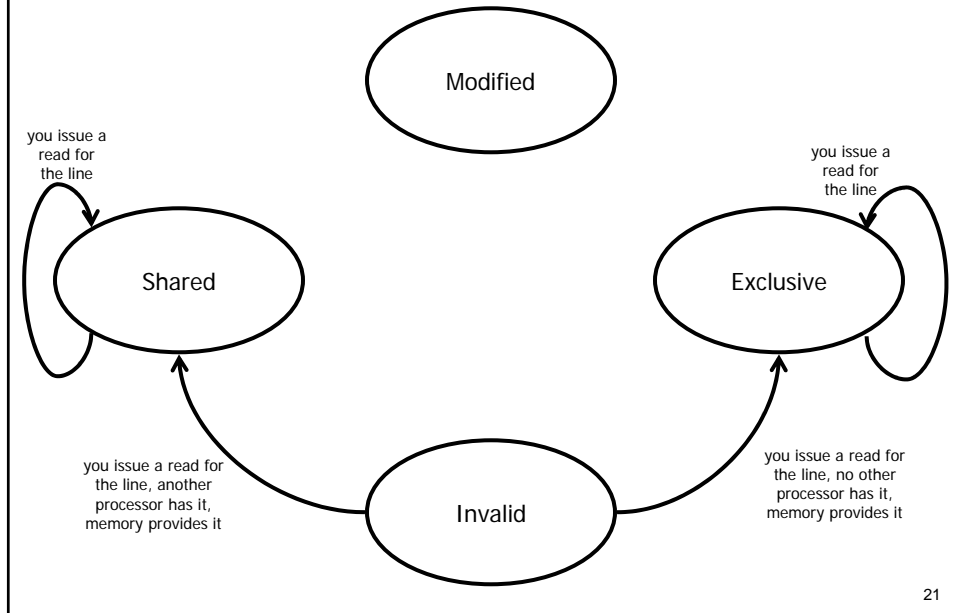
- A/B notation
 - if the cache controller observes the event A from the processor or the bus, then, in addition to state change, it generates the bus transaction or action B
 - A/-- means no transaction is generated on bus in response to event A

19

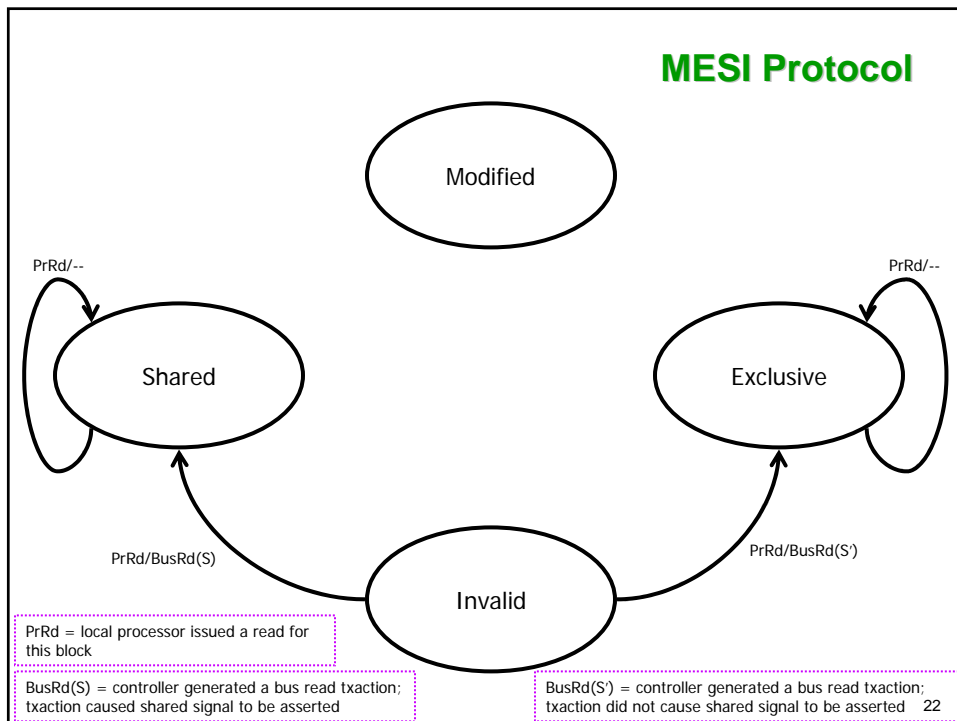
20

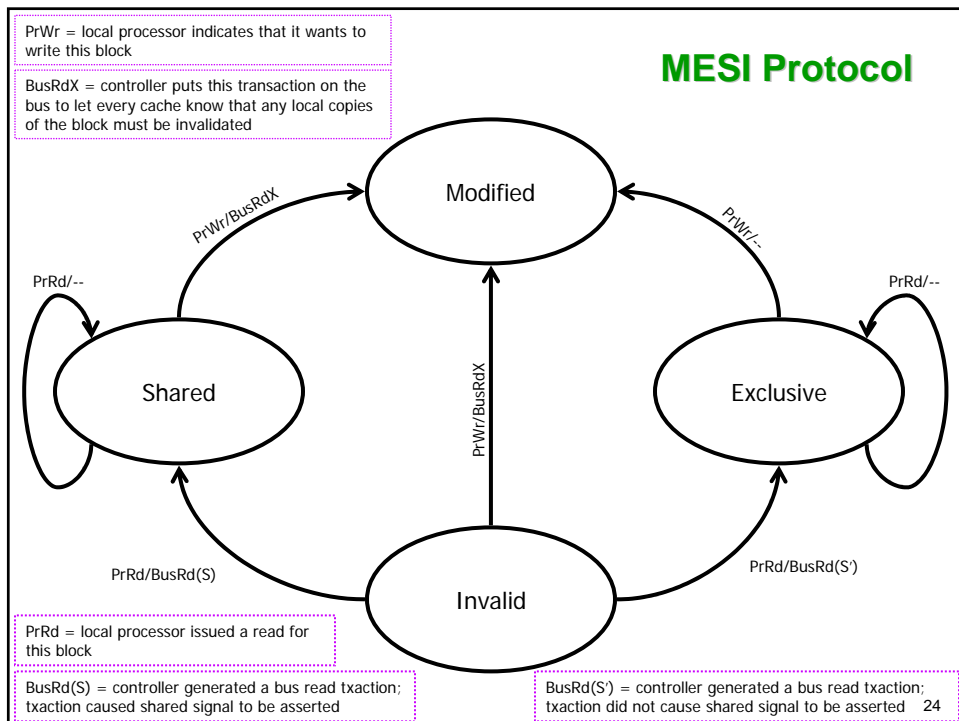
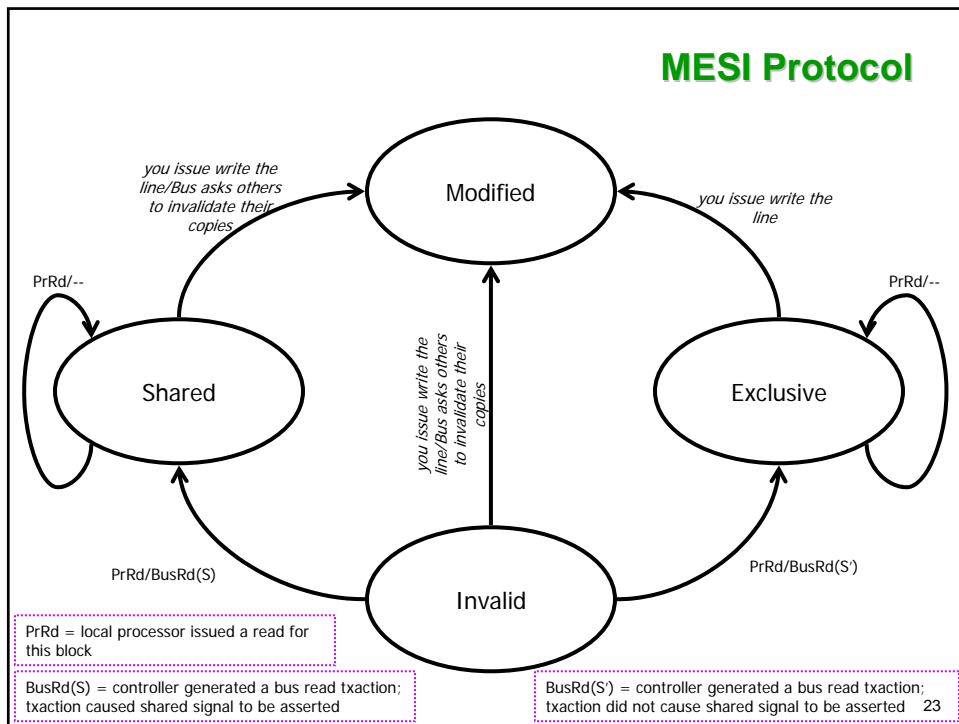
Repeated Slide

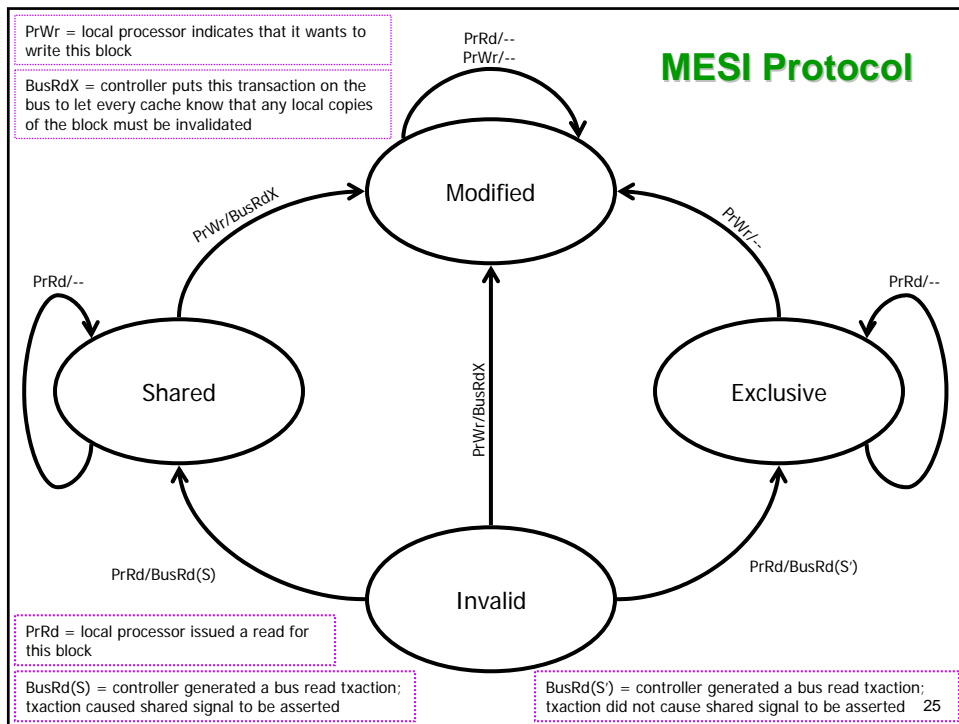
MESI Protocol



MESI Protocol



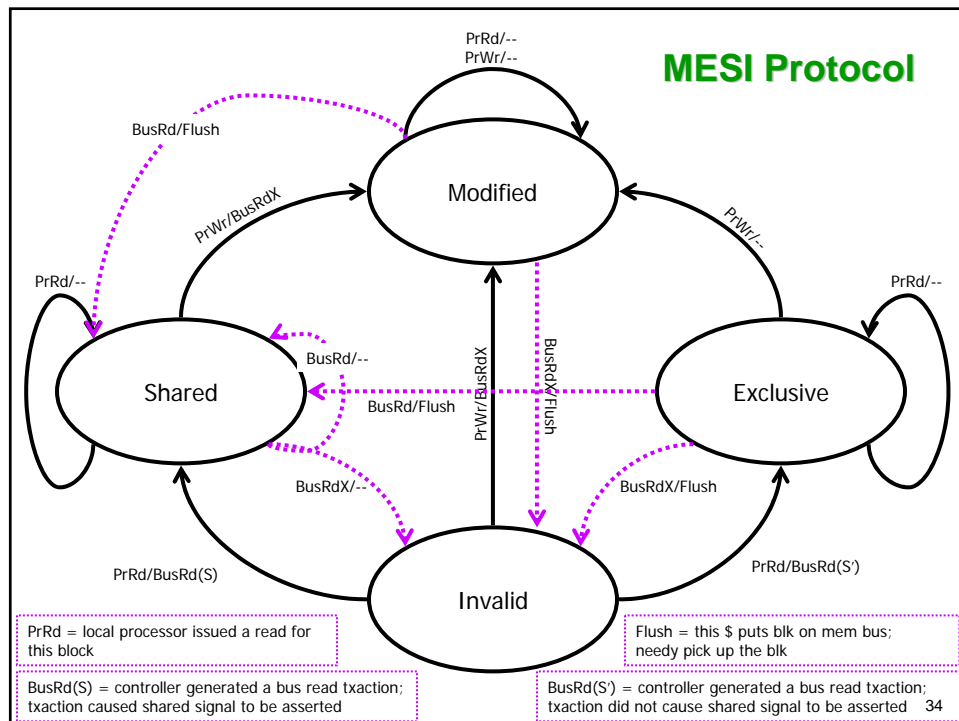




Alas!

- there is no line from shared to exclusive

33



Summary of MESI Protocol

- each cache watches all of the transactions on the memory bus
- if another processor requests a copy of a line in your cache, then you handle the request by sending the line over the bus and adjusting the state in your cache appropriately
- owner provides data
 - owner = cache with blk in M or E state
 - otherwise, owner = memory
- if your copy is S, E, or M, it becomes invalid if some other processor wants to write the data
- if your copy is S, E, or M, it becomes shared if some other processor wants to read the data

35